

CS/ECE/ME532 Assignment 10

1. Neural net functions

- a) Sketch the function generated by the following 3-neuron ReLU neural network.


$$f(x) = 2(x - 0.5)_+ - 2(2x - 1)_+ + 4(0.5x - 2)_+$$

where $x \in \mathbb{R}$ and where $(z)_+ = \max(0, z)$ for any $z \in \mathbb{R}$. Note that this is a single-input, single-output function. Plot $f(x)$ vs x by hand.

- b) Consider the continuous function depicted below. Approximate this function with ReLU neural network with 2 neurons. The function should be in the form

$$f(x) = \sum_{j=1}^2 v_j (w_j x + b_j)_+$$

Indicate the weights and biases of each neuron and sketch the neural network function.



- c) A neural network f_w can be used for binary classification by predicting the label as $\hat{y} = \text{sign}(f_w(\mathbf{x}))$. Consider a setting where $\mathbf{x} \in \mathbb{R}^2$ and the desired classifier is -1 if both elements of \mathbf{x} are less than or equal to zero and $+1$ otherwise. Sketch the desired classification regions in the two-dimensional plane, and provide a formula for a ReLU network with 2-neurons that can produce the desired classification. For simplicity, assume in this questions that $\text{sign}(0) = -1$.

2. **Gradients of a neural net.** Consider a 2 layer neural network of the form $f(\mathbf{x}) = \sum_{j=1}^J v_j (\mathbf{w}_j^T \mathbf{x})_+$. Suppose we want to train our network on a dataset of N samples \mathbf{x}_i with corresponding labels y_i , using a least squares loss function $\mathcal{L} = \sum_{i=1}^n (f(\mathbf{x}_i) - y_i)^2$. Derive the gradient descent update steps for the input weights \mathbf{w}_j and output weights v_j .

- 3. Compressing neural nets.** Large neural network models can be approximated by considering low rank approximations to weight matrices. The neural network $f(\mathbf{x}) = \sum_{j=1}^J \mathbf{v}_j(\mathbf{w}_j^T \mathbf{x})_+$ can be written as

$$f(\mathbf{x}) = \mathbf{v}^T(\mathbf{W}\mathbf{x})_+.$$

where \mathbf{v} is a $J \times 1$ vector of the output weights and \mathbf{W} is a $J \times d$ matrix with i th row \mathbf{w}_j^T . Let $\sigma_1, \sigma_2, \dots$ denote the singular values of \mathbf{W} and assume that $\sigma_i \leq \epsilon$ for $i > r$. Let f_r denote the neural network obtained by replacing \mathbf{W} with its best rank r approximation $\hat{\mathbf{W}}_r$. Assuming that \mathbf{x} has unit norm, find an upper bound to the difference $\max_{\mathbf{x}} |f(\mathbf{x}) - f_r(\mathbf{x})|$. (Hint: for any pair of vectors \mathbf{a} and \mathbf{b} , the following inequality holds $\|\mathbf{a}_+ - \mathbf{b}_+\|_2 \leq \|\mathbf{a} - \mathbf{b}\|_2$).

- 4. Face Emotion Classification with a three layer neural network.** In this problem we return to the face emotion data studied previously. You may find it very helpful to use code from an activity (or libraries such as Keras and Tensorflow).

- a) Build a classifier using a full connected three layer neural network with logistic activation functions. Your network should
- take a vector $\mathbf{x} \in \mathbb{R}^{10}$ as input (nine features plus a constant offset),
 - have a single, fully connected hidden layer with 32 neurons
 - output a scalar \hat{y} .

Note that since the logistic activation function is always positive, your decision should be as follows: $\hat{y} > 0.5$ corresponds to a ‘happy’ face, while $\hat{y} \leq 0.5$ is not happy.

- b) Train your classifier using stochastic gradient descent (start with a step size of $\alpha = 0.05$) and create a plot with the number of epochs on the horizontal axis, and training accuracy on the vertical axis. Does your classifier achieve 0% training error? If so, how many epoch does it take for your classifier to achieve perfect classification on the training set?
- c) Find a more realistic estimate of the accuracy of your classifier by using 8-fold cross validation. Can you achieve perfect test accuracy?

S32
Assignment 10
DEVIN
BREMER


1. Neural net functions

a) Sketch the function generated by the following 3-neuron ReLU neural network.

$$f(x) = 2(x - 0.5)_+ - 2(2x - 1)_+ + 4(0.5x - 2)_+$$

where $x \in \mathbb{R}$ and where $(z)_+ = \max(0, z)$ for any $z \in \mathbb{R}$. Note that this is a single-input, single-output function. Plot $f(x)$ vs x by hand.


1a.)



b) Consider the continuous function depicted below. Approximate this function with ReLU neural network with 2 neurons. The function should be in the form

$$f(x) = \sum_{j=1}^2 v_j (w_j x + b_j)_+$$

Indicate the weights and biases of each neuron and sketch the neural network function.



(b.)

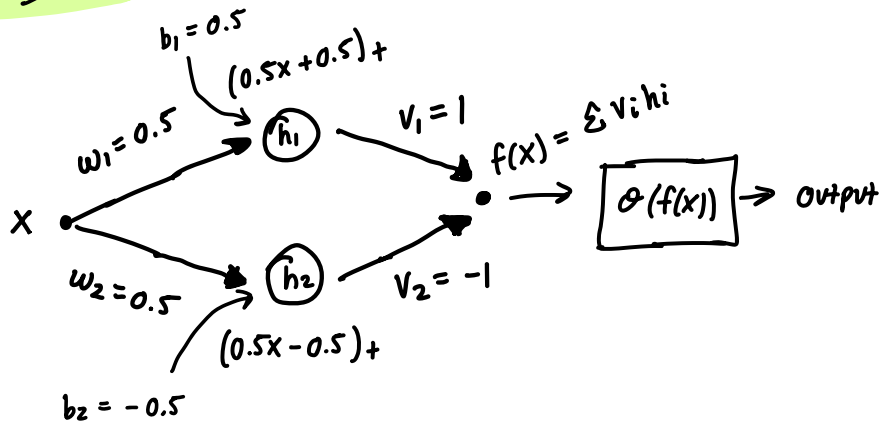
$$f(x) = (0.5x + 0.5)_+ - (0.5x - 0.5)_+$$

$$= \sum_{j=1}^2 v_j (w_j x + b_j)_+$$

$$\underline{v} = \begin{bmatrix} 1 \\ -1 \end{bmatrix}, \quad \underline{w} = \begin{bmatrix} 0.5 \\ 0.5 \end{bmatrix}$$


$$\underline{b} = \begin{bmatrix} 0.5 \\ -0.5 \end{bmatrix}$$

Neural Network:



- c) A neural network f_w can be used for binary classification by predicting the label as $\hat{y} = \text{sign}(f_w(\mathbf{x}))$. Consider a setting where $\mathbf{x} \in \mathbb{R}^2$ and the desired classifier is -1 if both elements of \mathbf{x} are less than or equal to zero and $+1$ otherwise. Sketch the desired classification regions in the two-dimensional plane, and provide a formula for a ReLU network with 2-neurons that can produce the desired classification. For simplicity, assume in this questions that $\text{sign}(0) = -1$.

$$\hat{y} = \text{sign}(f_w(\underline{x})). \quad \hat{y} = \begin{cases} -1, & x_1 \leq 0 \text{ \& \& } x_2 \leq 0 \\ 1, & \text{otherwise} \end{cases}$$



$$\text{blue circle} \rightarrow \hat{y} = -1$$

$$\text{green circle} \rightarrow \hat{y} = 1$$

$$\text{red line} \rightarrow \text{ReLU function}$$

$$f_w(\underline{x}) = (x_1)_+ + (x_2)_+$$

$$\hat{y} = \text{sign}(f_w(\underline{x})) = \begin{cases} -1, & x_1 \leq 0 \text{ \& \& } x_2 \leq 0 \\ 1, & \text{otherwise} \end{cases}$$

(assuming $\text{sign}(0) = -1$)

2. **Gradients of a neural net.** Consider a 2 layer neural network of the form $f(\mathbf{x}) = \sum_{j=1}^J v_j (\mathbf{w}_j^T \mathbf{x})_+$. Suppose we want to train our network on a dataset of N samples \mathbf{x}_i with corresponding labels y_i , using a least squares loss function $\mathcal{L} = \sum_{i=1}^N (f(\mathbf{x}_i) - y_i)^2$. Derive the gradient descent update steps for the input weights \mathbf{w}_j and output weights v_j .

$$\underline{v_j}$$

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial v_j} &= \sum_{i=1}^N 2(f(\underline{x}_i) - y_i) \cdot \frac{\partial f(\underline{x}_i)}{\partial v_j} \\ &= \sum_{i=1}^N (2(f(\underline{x}_i) - y_i) \cdot \frac{\partial \sum_{j=1}^J v_j (\underline{w}_j^T \underline{x})_+}{\partial v_j}) \\ &= \sum_{j=1}^N (2(f(\underline{x}_i) - y_i) \cdot (\underline{w}_j^T \underline{x})_+) \end{aligned}$$

w_j

$$\frac{\partial L}{\partial w_j} = \sum_{i=1}^N 2(f(x_i) - y_i) \cdot \frac{\partial f(x_i)}{\partial w_j}$$

$$= \sum_{i=1}^N 2(f(x_i) - y_i) \cdot \frac{\partial \sum_{j=1}^J v_j (w_j^T x_i)}{\partial w_j}$$

$\frac{d}{dx} \text{ReLU}(x) = \mathbb{1}_{\{x > 0\}}$

$$= \sum_{i=1}^N 2(f(x_i) - y_i) \cdot v_j \cdot \mathbb{1}_{\{w_j^T x_i > 0\}} \cdot x_i$$

Gradient Descent

1) Initialize $w_{m,j}^{(0)}, v_{k,l}^{(0)}$

2) For $t=0, 1, 2, \dots$

• Choose $i_t \in \{0, 1, \dots, N\}$ at random

• Compute $h_m^{i_t}, d_q^{i_t}$ from $x_{i_t}, w_{m,j}^t, v_{k,l}^t$

V update
Step

$$\rightarrow v_{k,l}^{(t+1)} = v_{k,l}^t - \alpha_t \cdot \sum_{j=1}^N 2(f(x_j) - y_j) \cdot (w_j^T x)$$

w update
Step

$$\rightarrow w_{m,j}^{(t+1)} = w_{m,j}^t - \alpha_t \cdot \sum_{i=1}^N 2(f(x_i) - y_i) \cdot v_j \cdot \mathbb{1}_{\{w_j^T x_i > 0\}} \cdot x_i$$

3. **Compressing neural nets.** Large neural network models can be approximated by considering low rank approximations to weight matrices. The neural network $f(\mathbf{x}) = \sum_{j=1}^J v_j (\mathbf{w}_j^T \mathbf{x})_+$ can be written as

$$f(\mathbf{x}) = \mathbf{v}^T (\mathbf{W} \mathbf{x})_+,$$

where \mathbf{v} is a $J \times 1$ vector of the output weights and \mathbf{W} is a $J \times d$ matrix with i th row \mathbf{w}_i^T . Let $\sigma_1, \sigma_2, \dots$ denote the singular values of \mathbf{W} and assume that $\sigma_i \leq \epsilon$ for $i > r$. Let f_r denote the neural network obtained by replacing \mathbf{W} with its best rank r approximation $\hat{\mathbf{W}}_r$. Assuming that \mathbf{x} has unit norm, find an upper bound to the difference $\max_{\mathbf{x}} |f(\mathbf{x}) - f_r(\mathbf{x})|$. (Hint: for any pair of vectors \mathbf{a} and \mathbf{b} , the following inequality holds $\|\mathbf{a}_+ - \mathbf{b}_+\|_2 \leq \|\mathbf{a} - \mathbf{b}\|_2$).

$$f(\underline{x}) = \sum_{j=1}^J v_j (\underline{w}_j^T \underline{x})_+ = \underline{v}^T (\underline{W} \underline{x})_+$$

$$f_r(\underline{x}) = \underline{v}^T (\hat{\underline{W}} \underline{x})_+$$

$$\rightarrow |f(\underline{x}) - f_r(\underline{x})| = \left\| (\underline{v}^T \underline{W} \underline{x})_+ - (\underline{v}^T \hat{\underline{W}} \underline{x})_+ \right\|_2 = \|\underline{v}\| \left\| (\underline{W} \underline{x})_+ - (\hat{\underline{W}} \underline{x})_+ \right\|$$

$$\rightarrow \|\underline{v}\|_2 \left\| (\underline{W} \underline{x})_+ - (\hat{\underline{W}} \underline{x})_+ \right\|_2 \leq \|\underline{v}\|_2 \underbrace{\left\| \underline{W} \underline{x} - \hat{\underline{W}} \underline{x} \right\|_2}_{\text{all } \delta_i \text{'s} < \epsilon \text{ for } i > r}$$

Thus this error can be at most ϵ since \underline{x} is unit norm

$$\rightarrow \|\underline{v}\|_2 \left\| (\underline{W} \underline{x})_+ - (\hat{\underline{W}} \underline{x})_+ \right\|_2 \leq \|\underline{v}\|_2 \cdot \epsilon$$

$$\rightarrow |f(\underline{x}) - f_r(\underline{x})| \leq \|\underline{v}\|_2 \cdot \epsilon$$

```

import numpy as np
import matplotlib.pyplot as plt
from scipy.io import loadmat
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

dataset = loadmat('face_emotion_data.mat')

X, y = dataset['X'], dataset['y']
n, p = np.shape(X)

y[y==1] = 0 # use 0/1 for labels instead of -1/+1
X = np.hstack((np.ones((n,1)), X)) # append a column of ones

```

4a

```

# Problem 4a - Devin Bresser

# define the model as follows:
# takes in a vector  $x \in \mathbb{R}^{10}$ 
# one hidden layer with 32 neurons
# outputs a scalar  $y_{\text{hat}} \in \mathbb{R}$ 

model = Sequential([
    Dense(32, activation='sigmoid', input_shape=(10,)), # hidden layer 32 neurons
    Dense(1, activation='sigmoid') # output layer
])

# compile the model with squared error loss function and SGD optimizer with learning rate 0.05
model.compile(optimizer=tf.keras.optimizers.SGD(learning_rate=0.05), loss='mean_squared_error', metrics=['accuracy'])

# Problem 4b - training accuracy

# try with 30 epochs
max_epochs = 30
train_accuracies = []

for i in range(0, max_epochs):

    fit = model.fit(X, y, epochs=1, batch_size=1, verbose=0)

    # obtain the accuracy for this epoch
    train_accuracy = fit.history['accuracy'][0]
    train_accuracies.append(train_accuracy)

    #print(f"Completed Epochs: {i}, Accuracy: {train_accuracy:.4f}")

first_perfect_accuracy = next((i, acc) for i, acc in enumerate(train_accuracies, start=1) if acc == 1)

print(f"training accuracy of 1.0 first occurs at epoch # {first_perfect_accuracy[0]}")


# plotting:
epochs = list(range(1, max_epochs+1))

plt.plot(epochs, train_accuracies, marker='o', linestyle='--')
plt.title('Training Accuracy over Epochs')
plt.xlabel('Epoch Number')
plt.ylabel('Training Accuracy')
plt.grid(True)
plt.show()

```

4b

training accuracy of 1.0 first occurs at epoch # 17



Yes, 100% training accuracy is achievable with this model!

4c

```

# Problem 4c - CV
# This code essentially runs the cross validation process with a number of epochs
# ranging from 1 to 100. We should see a convergence after a certain point,
# and from there be able to tell if it's possible to attain perfect test accuracy.

from sklearn.model_selection import KFold

max_epochs = 100
epoch_accuracies = []

# for each number of epochs, do the CV
for i in range(1, max_epochs):
    kf = KFold(n_splits=8, shuffle=True, random_state=42)
    fold_accuracies = []

    # do CV using KFold module
    for train_index, test_index in kf.split(X):

        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]

        # re-define the model each time
        model = Sequential([
            Dense(32, activation='sigmoid', input_shape=(10,)),
            Dense(1, activation='sigmoid')
        ])
        model.compile(optimizer=tf.keras.optimizers.SGD(learning_rate=0.05),
                      loss='mean_squared_error',
                      metrics=['accuracy'])

        # fit the model to the training part
        model.fit(X_train, y_train, epochs=i, batch_size=1, verbose=0) # picked 20 epochs based on training results

        # get the accuracy on the testing part and append to fold_accuracies
        _, accuracy = model.evaluate(X_test, y_test, verbose=0)
        fold_accuracies.append(accuracy)

    # compute the average accuracy across all folds
    average_accuracy = np.mean(fold_accuracies)

    # store average test accuracy for that epoch count
    print(f"average test accuracy at epoch {i}: {average_accuracy}")
    epoch_accuracies.append(average_accuracy)

```

```

average test accuracy at epoch 1: 0.7578125
average test accuracy at epoch 2: 0.8046875
average test accuracy at epoch 3: 0.8515625
average test accuracy at epoch 4: 0.84375
average test accuracy at epoch 5: 0.8828125
average test accuracy at epoch 6: 0.8828125
average test accuracy at epoch 7: 0.8984375
average test accuracy at epoch 8: 0.8828125
average test accuracy at epoch 9: 0.9140625
average test accuracy at epoch 10: 0.9375
average test accuracy at epoch 11: 0.9296875
average test accuracy at epoch 12: 0.9375
average test accuracy at epoch 13: 0.9296875
average test accuracy at epoch 14: 0.9375
average test accuracy at epoch 15: 0.9453125
average test accuracy at epoch 16: 0.953125
average test accuracy at epoch 17: 0.96875
average test accuracy at epoch 18: 0.9765625
average test accuracy at epoch 19: 0.9765625
average test accuracy at epoch 20: 0.9765625
average test accuracy at epoch 21: 0.9609375
average test accuracy at epoch 22: 0.9609375
average test accuracy at epoch 23: 0.96875
average test accuracy at epoch 24: 0.96875
average test accuracy at epoch 25: 0.9765625
average test accuracy at epoch 26: 0.9765625
average test accuracy at epoch 27: 0.9765625
average test accuracy at epoch 28: 0.96875
average test accuracy at epoch 29: 0.984375
average test accuracy at epoch 30: 0.984375
average test accuracy at epoch 31: 0.96875
average test accuracy at epoch 32: 0.96875
average test accuracy at epoch 33: 0.9765625
average test accuracy at epoch 34: 0.9921875
average test accuracy at epoch 35: 0.9765625
average test accuracy at epoch 36: 0.984375
average test accuracy at epoch 37: 0.984375
average test accuracy at epoch 38: 0.96875
average test accuracy at epoch 39: 0.9921875
average test accuracy at epoch 40: 0.984375
average test accuracy at epoch 41: 0.9765625
average test accuracy at epoch 42: 0.9765625
average test accuracy at epoch 43: 0.9921875
average test accuracy at epoch 44: 0.9765625
average test accuracy at epoch 45: 0.9765625
average test accuracy at epoch 46: 0.9765625
average test accuracy at epoch 47: 0.9765625
average test accuracy at epoch 48: 0.984375
average test accuracy at epoch 49: 0.984375
average test accuracy at epoch 50: 0.9765625
average test accuracy at epoch 51: 0.984375
average test accuracy at epoch 52: 0.96875

```

```

average test accuracy at epoch 48: 0.984375
average test accuracy at epoch 49: 0.984375
average test accuracy at epoch 50: 0.9765625
average test accuracy at epoch 51: 0.984375
average test accuracy at epoch 52: 0.96875
average test accuracy at epoch 53: 0.9765625
average test accuracy at epoch 54: 0.984375
average test accuracy at epoch 55: 0.9765625
average test accuracy at epoch 56: 0.984375
average test accuracy at epoch 57: 0.984375
average test accuracy at epoch 58: 0.984375
average test accuracy at epoch 59: 0.984375
average test accuracy at epoch 60: 0.9765625
average test accuracy at epoch 61: 0.984375
average test accuracy at epoch 62: 0.984375
average test accuracy at epoch 63: 0.9765625
average test accuracy at epoch 64: 0.9765625
average test accuracy at epoch 65: 0.984375
average test accuracy at epoch 66: 0.984375
average test accuracy at epoch 67: 0.984375
average test accuracy at epoch 68: 0.9765625
average test accuracy at epoch 69: 0.9765625
average test accuracy at epoch 70: 0.9765625
average test accuracy at epoch 71: 0.984375
average test accuracy at epoch 72: 0.9765625
average test accuracy at epoch 73: 0.9765625
average test accuracy at epoch 74: 0.984375
average test accuracy at epoch 75: 0.984375
average test accuracy at epoch 76: 0.984375
average test accuracy at epoch 77: 0.984375
average test accuracy at epoch 78: 0.9765625
average test accuracy at epoch 79: 0.9765625
average test accuracy at epoch 80: 0.9765625
average test accuracy at epoch 81: 0.9765625
average test accuracy at epoch 82: 0.9765625
average test accuracy at epoch 83: 0.984375
average test accuracy at epoch 84: 0.9765625
average test accuracy at epoch 85: 0.984375
average test accuracy at epoch 86: 0.9765625
average test accuracy at epoch 87: 0.984375
average test accuracy at epoch 88: 0.9765625
average test accuracy at epoch 89: 0.984375
average test accuracy at epoch 90: 0.9765625
average test accuracy at epoch 91: 0.984375
average test accuracy at epoch 92: 0.984375
average test accuracy at epoch 93: 0.984375
average test accuracy at epoch 94: 0.9765625
average test accuracy at epoch 95: 0.984375
average test accuracy at epoch 96: 0.984375
average test accuracy at epoch 97: 0.9765625
average test accuracy at epoch 98: 0.984375
average test accuracy at epoch 99: 0.984375

```

... Conclusion:

100% avg. test accuracy isn't

possible with this model.